Detecting Changes in Power Plant NO_x Emissions Using Deep Learning and TEMPO Satellite Imagery

Pranav Walimbe
University of California, Berkeley
Berkeley, California, USA
pranav.walimbe@berkeley.edu

Hikari Murayama
University of California, Berkeley
Berkeley, California, USA
hikari_murayama@berkeley.edu

Ron Cohen
University of California, Berkeley
Berkeley, California, USA
rccohen@berkeley.edu

Abstract

NASA's Tropospheric Emissions: Monitoring of Pollution (TEMPO) instrument provides hourly daytime observations of tropospheric nitrogen dioxide (NO_2) across North America. These observations could support independent, facility-level monitoring of power plant nitrogen oxides (NO_x) emissions, filling gaps when continuous emissions monitoring system (CEMS) measurements fail or facilities are not subject to hourly emissions reporting requirements. Previous studies have estimated hourly point-source emissions using empirical methods, which are highly sensitive to instrument retrieval uncertainty, changing atmospheric conditions, TEMPO's coarse spatial resolution, signals from nearby emission sources, and temporal mismatch with hourly CEMS records. We instead formulate facility monitoring as change detection over sequences of TEMPO observations. This allows the model to learn relative emissions behavior rather than estimate exact emissions from noisy data. We develop a convolutional gated recurrent unit (ConvGRU) model that processes five consecutive TEMPO observations together with meteorological inputs, features representing plume transport, and tabular features encoding time of day and seasonality to classify power plant emission trajectories as increasing, steady, or decreasing. On a held-out evaluation set of more than 22,000 samples covering 116 active U.S. power plants observed between August 2023 and June 2026, the ConvGRU model achieved a macro-F1 score of 0.56, outperforming a baseline multilayer perceptron (MLP) trained only on the same time-derived features, which scored 0.45. Our findings suggest that the ConvGRU model can identify signals of power plant emissions from TEMPO observations beyond the time-derived features shared with the baseline MLP. More broadly, they indicate that change detection may offer a more tractable framework for modeling power plant activity from TEMPO observations than estimating absolute emissions.

1 Introduction

Based on U.S. Environmental Protection Agency (EPA) inventory data, power plants in the U.S. electricity generation sector emitted approximately 0.74 million short tons of NO_x in 2023, about 9.5% of national emissions [30]. NO_x emissions contribute to ground-level ozone and secondary particulate matter, pollutants that can inflame and damage the airways, reduce lung function, and aggravate asthma [28]. EPA regulations require most large U.S. fossil-fuel-fired generating units that produce electricity for sale to submit unit-level hourly emissions data on a quarterly basis [33]. However, monitoring equipment malfunctions and maintenance can make CEMS measurements unavailable or invalid, requiring plants to report calculated substitute values instead of direct measurements [29]. Behind-the-meter generators that do not sell electricity can fall outside federal hourly reporting requirements, such as the natural gas turbines powering SpaceXAI's Colossus 2 data center [11]. Outside the United States, comparable public, hourly, unit-level emissions data are generally unavailable; global inventories typically provide only annual estimates [12]. Independent estimates could therefore fill reporting gaps and provide timely insight into facility operations ahead of the EPA's quarterly release of hourly emissions data.

TEMPO is an ultraviolet-visible spectrometer hosted aboard a geostationary satellite that observes tropospheric vertical column densities (TVCDs) for various trace gases across North America during daylight hours [37]. Earlier polar-orbiting instruments, including the Ozone Monitoring Instrument (OMI) and TROPOspheric Monitoring Instrument (TROPOMI), typically observe a given location once per day; TEMPO's geostationary view provides approximately hourly scans with finer spatial resolution [15, 35, 37]. Hourly sampling captures intraday changes in emissions plumes and atmospheric transport that daily observations can miss. However, power plant emission signals in TEMPO NO_2 TVCDs can be weakened or obscured by factors including evolving meteorology and atmospheric chemistry, cloud cover, viewing geometry and retrieval uncertainty, dilution within TEMPO's 2.0 km × 4.75 km spatial resolution, and interference from nearby emission sources [11, 12, 24]. Comparisons with CEMS records, our most reliable ground truth, introduce an additional temporal mismatch because TEMPO observes a facility at a specific time during each hourly scan, whereas CEMS data aggregate emissions over an entire clock hour [24, 33].

Satellite instruments measure atmospheric column abundance rather than emissions directly. Lightweight inversion methods therefore infer source rates from the spatial structure of an observed plume and estimated winds. Cross-sectional flux (CSF) removes the background column, integrates the remaining enhancement across transects perpendicular to the plume, and multiplies the resulting line density by an effective wind speed. Under steady-state conditions, this flux approximates the source emission rate. The directional derivative approach (DDA) instead combines the wind field with spatial gradients in the observed column to estimate a local source term, which is then integrated around the facility [20, 23].

Although promising, these methods remain sensitive to source characteristics and observing conditions. Kuhlmann et al. applied CSF to TROPOMI NO_2 observations of 18 power plants in Europe and the United States, recovering seasonal NO_x emissions variability. However, filtering for cloud cover and interference from nearby plumes left only 38–170 valid overpasses per plant per year, with particularly sparse winter coverage. These sampling gaps, together with uncertainty in chemistry corrections, limited the ability to

resolve short-term emissions fluctuations [14]. With TEMPO observations, Sun et al. improved correlations with CEMS at 10 of 14 power plants by incorporating changes between consecutive NO_2 columns, although the resulting estimates remained approximately fourfold lower than CEMS measurements [23]. Beaudry et al. used a model that separately fits NO_2 buildup near the source and plume decay downwind. Their average emissions estimates were within 50% of CEMS values at 27 of 64 plants, but bias was substantially greater below 250 kg NO_2 h^{-1} [3].

Machine learning (ML) studies have addressed several related point-source monitoring tasks. Huang and Wang trained a tabular MLP to predict whether a TROPOMI observation contained a detectable power plant NO_2 plume using 25 facility, meteorological, surface, and sensor variables sampled at the plant location [12]. Their models achieved F1 scores above 0.66, but detectability ranged from below 20% to above 60% for plants with similar emissions. Their interference filters also retained plants representing only 45.0% of U.S. and 21.1% of global NO_x emissions in their datasets. Alnaim et al. used stacked long short-term memory (LSTM) networks to predict next-day EPA-reported NO_2 emissions from seven-day sequences of previous emissions, scalar TROPOMI observations, MERRA-2 meteorology, aerosol measurements, and calendar variables [1]. Their best single-plant LSTM had an error near 40%, while multi-plant experiments exceeded 55% error and often underestimated extreme values. Finch et al. developed a convolutional neural network (CNN) using 6,086 manually labeled TROPOMI images, achieving greater than 90% test accuracy for NO_2 plume detection. However, reliance on manual annotation limits the scalability of label generation, and their evaluation addresses plume presence rather than temporal changes in emissions [10].

The point-source ML studies reviewed above do not jointly model the spatial structure and temporal evolution of consecutive trace-gas satellite rasters. We instead formulate facility monitoring as change detection over five TEMPO scans. Classifying emissions as increasing, steady, or decreasing preserves information about short-term facility behavior while reducing reliance on recovering exact emissions amounts from noisy observations.

Two relevant model families are convolutional recurrent neural networks (CRNNs) and spatiotemporal vision transformers (ViTs). At each timestep, a CRNN combines features from the current raster with a spatially structured state carried from the preceding timestep. Convolutional operations within the recurrent cell preserve the raster grid and update each location using its surrounding neighborhood. ConvLSTM controls separate cell and hidden states through input, forget, and output gates, providing explicit control over memory retention at the cost of additional parameters and computation. ConvGRU instead maintains one hidden state controlled by reset and update gates, producing a more compact recurrent model with less explicit control over long-term memory [2, 22]. Spatiotemporal ViTs divide each raster into patches and use self-attention to connect patch representations across locations and timesteps. This provides direct access to global spatiotemporal relationships without recurrent processing, but supplies a weaker locality bias and shifts more pattern learning to the training data; the computational cost also grows with the number of patches and timesteps [4].

We select ConvGRU to balance modeling capacity with parameter efficiency. Each sample spans only five scans, so ConvLSTM's additional memory state offers less apparent benefit than it would for longer sequences while increasing the number of model parameters. The input rasters are also small (24×24 pixels), and power plant plumes generally form connected clusters of pixels whose evolution can be captured through local convolutional operations rather than long-range interactions. TEMPO also provides only about three years of observations, which limits the overall dataset size. We therefore use a compact ConvGRU to capture spatiotemporal structure while limiting model complexity.

2 Data Sources

2.1 TEMPO NO_2 Data

NASA distributes TEMPO NO_2 retrievals in two formats: Level 2 data on the instrument's native measurement footprints and Level 3 data regridded from each east-to-west scan [24]. We downloaded version 4 Level 2 NO_2 retrievals from August 2023 through June 2026. Along with NO_2 TVCDs, the product provides metadata including native footprint geometry, mirror-step observation times, main data quality flags, and effective cloud fractions, which we later use for regridding, quality filtering, and temporal alignment with CEMS records.

2.2 CEMS Records

From the EPA's Clean Air Markets Program Data (CAMPD) APIs, we downloaded hourly emissions records for individual U.S. power plant units from August 2023 through June 2026 and annual facility and unit attributes [34]. From the hourly records, we use facility and unit identifiers, reporting date and hour, NO_x mass and its measurement flag, and operating time. From the annual attributes, we use facility coordinates.

2.3 HRRR Weather Data

The High-Resolution Rapid Refresh (HRRR) is a weather prediction model updated hourly that covers the contiguous United States at a spatial resolution of 3 km [17]. From NOAA's public HRRR archive, we downloaded hourly 2 m air temperature and eastward and northward wind components at 80 m above ground for August 2023 through June 2026 [18].

3 Methods

3.1 AOI Construction

For each power plant, we define a 72 km $\times$ 72 km area of interest (AOI) centered on the facility and divide it into a 24×24 grid of 3 km $\times$ 3 km cells. For each AOI-hour, we sum the unit-level NO_x mass reported by all facilities in the AOI and discard the hour if any contributing measurement is invalid or unavailable. We score each AOI using its median total NO_x mass across hours whose mean unit operating time meets or exceeds the AOI-specific median. We retain the highest-scoring 50% to focus the dataset on emissions signals more likely to be detected by TEMPO.

3.2 TEMPO Regridding and Quality Filtering

TEMPO Level 2 observations lie on irregular native footprints, whereas our model uses a fixed spatial grid. We project the footprint boundaries and AOI grid to the contiguous-U.S. Albers equal-area

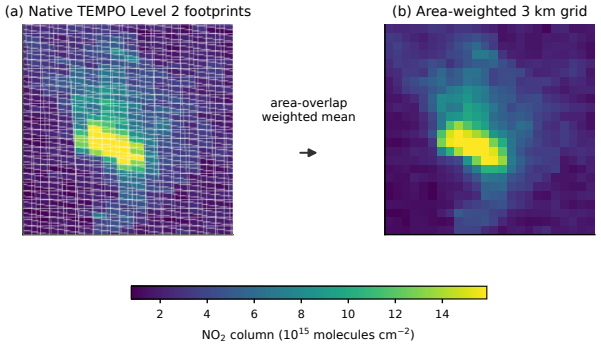


Figure 1: TEMPO NO_2 regridding within a 72 km × 72 km area centered on a power plant. (a) Native Level 2 footprints overlaid with the target grid. (b) Footprint values averaged by area of overlap onto a 24 × 24 grid of 3 km × 3 km cells.

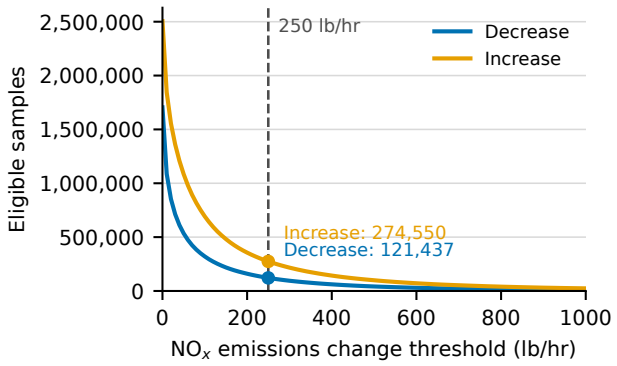


Figure 2: Eligible increase and decrease sample counts across NO_x emissions change thresholds, before raster quality filtering and class balancing.

coordinate system. We then assign each AOI grid cell an NO_2 TVCD by averaging the TEMPO footprints that intersect it, weighted by their area of overlap. Figure 1 illustrates this mapping from native footprints to the fixed AOI grid.

A footprint contributes to the NO_2 average only when its main data quality flag equals zero, indicating a reliable retrieval, and its effective cloud fraction does not exceed 0.20. Grid cells without an accepted footprint overlap remain masked, and we retain a binary mask of the original valid cells. We use bilinear interpolation to fill masked NO_2 pixels for downstream modeling.

3.3 Sample Construction

Each sample contains five consecutive observations of one AOI, indexed from t_0 to t_4. We retain a sequence only when each pair of adjacent TEMPO scans is separated by 40–70 minutes and every NO_2 raster has valid pixel coverage of at least 85%. At each timestep, we pair the TEMPO raster with the HRRR temperature and wind rasters from the closest hour and interpolate them to the same 24 × 24 raster size and spatial coverage as the NO_2 rasters.

Because CEMS records correspond to clock hours rather than TEMPO scan times, we linearly interpolate the AOI-hour NO_x emissions between the surrounding hourly CEMS records at each scan time. Let E_i denote the interpolated emissions at time t_i. For scans t_a through t_b, we calculate the time-weighted mean

$$\overline{E}_{a:b} = \frac{\sum_{i=a}^{b-1}(t_{i+1} - t_i)(E_i + E_{i+1})/2}{t_b - t_a}. \tag{1}$$

We define the sequence change as $\Delta E = \overline{E}_{2:4} - \overline{E}_{0:2}$ and use the threshold $\theta = 250\ \mathrm{lb\,h^{-1}}$ to assign its category:

$$c = \begin{cases} \text{increase}, & \Delta E \geq \theta, \\ \text{steady}, & -\theta < \Delta E < \theta, \\ \text{decrease}, & \Delta E \leq -\theta. \end{cases} \tag{2}$$

We chose the threshold $\theta = 250\ \mathrm{lb\,h^{-1}}$ based on the balance between valid sample counts across classes and separability between classes from larger emissions deltas, which we visualize in Figure 2.

Requiring a substantial emissions change helps make the classes distinguishable despite noise in the TEMPO NO_2 rasters. Even at the 250 $\mathrm{lb\,h^{-1}}$ threshold, however, the visual evidence can vary markedly between samples, as the two increase-labeled sequences in Figure 3 illustrate.

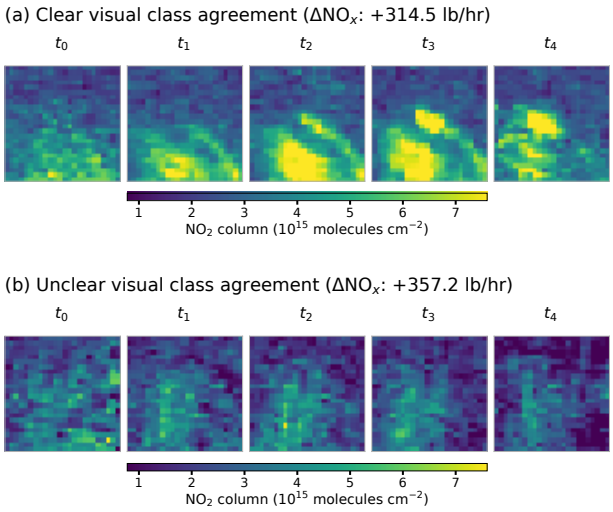


Figure 3: Two five-observation TEMPO NO_2 sequences with CEMS-derived increase labels. (a) An AOI centered on the Victor H. Braunig Power Station in San Antonio, Texas, observed on August 25, 2023, shows a clear visual pattern of increasing emissions. (b) An AOI centered on the High Bridge Generating Station in St. Paul, Minnesota, observed on September 16, 2025, shows little visual evidence of increasing emissions. Both sequences exceed the 250 $\mathrm{lb\,h^{-1}}$ NO_x emissions change threshold.

Figure 4 then shows a complete five-observation training sample, combining its NO_2 sequence with the other model inputs and its derived label.

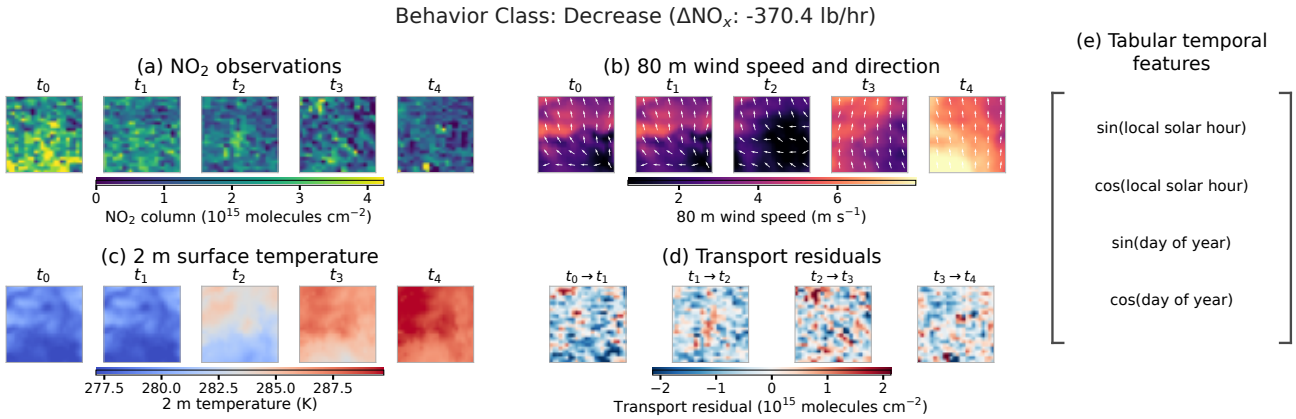


Figure 4: Five-observation training sample (validity masks excluded) from an AOI centered on the Hugh L. Spurlock Generating Station in Maysville, Kentucky, observed on February 26, 2025. Panels (a)–(c) show the NO_2, wind, and temperature inputs at t_0–t_4; (d) shows the four residuals between adjacent observations; and (e) shows the temporal features.

3.4 Transport Model

Emissions change signals in consecutive NO_2 rasters can be degraded by movement and decay of the existing plume. For each pair of consecutive NO_2 rasters, we use a simple transport model to predict the next observation without new emissions in the AOI. We then subtract this prediction from the observed next raster to obtain a transport residual, which separates plume structure explained by modeled advection, horizontal dispersion, and chemical decay from remaining variation that may reflect emissions changes within the AOI.

For the interval $\Delta t_i = t_{i+1} - t_i$, let $N_i(\boldsymbol{x})$ and $N_{i+1}(\boldsymbol{x})$ denote the earlier and later NO_2 rasters, $\boldsymbol{u}_i(\boldsymbol{x})$ and $\boldsymbol{u}_{i+1}(\boldsymbol{x})$ their wind fields, and B_i the median NO_2 value along the boundary of the earlier raster. Using these quantities, we construct the transported prediction and residual as follows:

$$\text{(i)} \quad \overline{\boldsymbol{u}}_i = \gamma[(1-\lambda)\boldsymbol{u}_i + \lambda \boldsymbol{u}_{i+1}]$$

$$\text{(ii)} \quad A_i(\boldsymbol{x}) = N_i(\boldsymbol{x} - \overline{\boldsymbol{u}}_i(\boldsymbol{x})\Delta t_i)$$

$$\text{(iii)} \quad \sigma_i = \sigma_0 \sqrt{\frac{\Delta t_i}{1\ \text{h}}}$$

$$\text{(iv)} \quad D_i = B_i + G_{\sigma_i} * (A_i - B_i)$$

$$\text{(v)} \quad \widehat{N}_{i+1} = B_i + \exp\left(-\frac{\Delta t_i}{\tau}\right)(D_i - B_i)$$

$$\text{(vi)} \quad R_i = N_{i+1} - \widehat{N}_{i+1}.$$

We selected the values of λ, γ, σ_0, and τ using the tuning procedure described in Section 3.8. Equations (i) and (ii) represent wind-driven advection. Satellite columns do not resolve the vertical distribution of plume mass, so point-source methods commonly use an effective wind to approximate its mean transport speed [16, 20]. We approximate this wind by averaging the fields at the two endpoints ($\lambda = 0.5$) and applying a modest speed correction ($\gamma = 0.8$). Equation (ii) then applies semi-Lagrangian advection by tracing each destination cell backward along the effective wind over Δt_i and bilinearly sampling the earlier raster at the resulting departure location.

Equations (iii) and (iv) approximate unresolved turbulent horizontal dispersion. Gaussian plume representations are widely used in satellite point-source studies [16, 20]. We set σ_0 to one 3 km grid cell after one hour and increase the width with $\sqrt{\Delta t_i}$. Applying the Gaussian kernel to $A_i - B_i$ spreads the plume enhancement while preserving the estimated background.

Equation (v) applies first-order loss to the above-background enhancement. The effective timescale $\tau = 8$ h represents the combined persistence of the observed NO_2 plume under chemical removal and other unresolved losses. Exponential loss is commonly used to describe the downwind decay of satellite-observed NO_x plumes, whose effective lifetimes vary with meteorology and plume chemistry [16]. Equation (vi) subtracts the resulting transport-only prediction from the next observation.

Each five-observation sample produces four transport-residual rasters, one for each pair of consecutive NO_2 rasters. For each pair, we advect and diffuse the earlier observation's validity mask alongside its NO_2 raster, then multiply it by the later observation's mask to form the residual's quality mask. Figure 4 shows all four residuals for the example sample.

3.5 Dataset Stratification

Our split design has three goals: prevent leakage from overlapping AOIs, produce an approximately 70%/15%/15% stratification of samples, and construct class distributions that support balanced training and comparable evaluation across the held-out splits. We group overlapping AOIs into connected clusters and treat each cluster as an indivisible assignment unit. Because each cluster contains an arbitrary number of AOIs, samples, and class instances, cluster-level assignment cannot generally meet every target exactly.

We use mixed-integer linear programming (MILP) to assign each cluster to exactly one split while minimizing normalized deviations from target counts of AOIs and eligible samples. The first optimization targets a 70%/15%/15% sample stratification using AOI and

sample counts, after which we fix the training clusters. A second optimization assigns the remaining clusters equally between validation and test while balancing AOI, sample, and per-class counts. It also constrains both held-out splits to have the same median number of plant units per AOI as the combined held-out pool.

We then uniformly downsample each training class to the size of the smallest class. In validation and test, we retain every increase and decrease sample and uniformly subsample steady samples to at most the size of the larger non-steady class.

We then apply the raster filtering described in Section 3.3. Across the three splits, 19.4% of selected decrease samples, 25.5% of steady samples, and 32.4% of increase samples are retained through raster filtering. The retention pattern is primarily driven by seasonal cloud cover: decrease samples occur more often in winter, steady samples are less seasonally concentrated, and increase samples occur more often in summer. Table 1 reports the final dataset and class counts.

Table 1: Final modeling dataset composition by split and behavior class. AOIs are distinct 72 km × 72 km bounding boxes centered on power plants. An active plant has nonzero operating time during at least one observed hour in the dataset.

Split	AOIs	Active plants	Decrease	Steady	Increase	Total
Training	469	530	17,130	23,289	29,888	70,307
Validation	100	109	3,240	7,648	9,902	20,790
Test	98	116	3,142	8,722	10,602	22,466
Total	667	755	23,512	39,659	50,392	113,563

3.6 ConvGRU Model

Each sample supplies five 24×24 observation frames and four transport-residual frames, one for each consecutive NO_2 transition. An observation frame contains NO_2, its quality mask, eastward and northward 80 m winds, and 2 m temperature. A residual frame contains the transport residual and its quality mask. The model also receives four time-derived tabular features computed at the sequence midpoint.

Separate convolutional encoders process the observation and residual frames, sharing weights across time within each branch. A ConvGRU updates a spatial hidden state over four transitions. The model then combines the final visual representation with the time-derived features to classify emissions change behavior. Figure 5 shows the architecture and one ConvGRU cell.

We denote the observation and residual encoders E_x and E_r, respectively, in Figure 5(a). Separate encoders account for the different meanings of their inputs. Both encoders begin with a 5×5 convolution that produces 32 channels, followed by five residual blocks. Strided blocks increase the channel width from 32 to 48 and then 64 while reducing each 24×24 raster to a $64 \times 6 \times 6$ feature map. Each residual block applies two 3×3 convolutions with group normalization, a SiLU activation after the first convolution, and another SiLU activation after adding the skip connection. Blocks that change channel width or spatial resolution use a 1×1 convolution on the skip path.

As shown in Figure 5(a), the encoded t_0 observation initializes the hidden state, so the model does not synthesize a residual before the first observation. At each subsequent timestep t_i, the model concatenates the encoded observation with the encoded residual for $t_{i-1} \rightarrow t_i$. A 1×1 projection fuses the two encoded feature maps before the recurrent update.

Figure 5(b) expands one recurrent update. Let $\boldsymbol{f}_i$ denote the fused input and $\boldsymbol{h}_{i-1}$ the previous hidden state, a $64 \times 6 \times 6$ feature map that carries spatial information from earlier observations. The ConvGRU computes

$$[\boldsymbol{r}_i, \boldsymbol{z}_i] = \sigma\big(W_g * [\boldsymbol{f}_i, \boldsymbol{h}_{i-1}] + \boldsymbol{b}_g\big), \tag{3}$$

$$\widetilde{\boldsymbol{h}}_i = \tanh\big(W_h * [\boldsymbol{f}_i, \boldsymbol{r}_i \odot \boldsymbol{h}_{i-1}] + \boldsymbol{b}_h\big), \tag{4}$$

$$\boldsymbol{h}_i = (1 - \boldsymbol{z}_i) \odot \boldsymbol{h}_{i-1} + \boldsymbol{z}_i \odot \widetilde{\boldsymbol{h}}_i, \tag{5}$$

where the subscripts g and h denote the gate and candidate-state convolutions, respectively. W_g and W_h are their learned 3×3 kernels, while $\boldsymbol{b}_g$ and $\boldsymbol{b}_h$ are the corresponding biases. The symbol $*$ denotes convolution, and $[\cdot, \cdot]$ denotes channel-wise concatenation. The reset gate $\boldsymbol{r}_i$ controls how much previous spatial information contributes to the candidate state, while the update gate $\boldsymbol{z}_i$ determines how much of that candidate replaces the previous state.

After the fourth recurrent update, global average and maximum pooling summarize distributed and localized evidence, respectively, in the final hidden state. The model concatenates their outputs and projects them into a visual embedding. A feedforward encoder maps four time-derived tabular features (the sine and cosine of local solar hour and of day of year) into a tabular embedding. The model adds this embedding after recurrent processing, keeping the tabular features out of the spatial state updates. The fused representation passes through a classification head to produce logits for the three behavior classes.

3.7 Baseline MLP Model

We use a lightweight baseline MLP to gauge how much emissions change behavior can be predicted from recurring daily and annual patterns without satellite or weather inputs. It uses the same four time-derived tabular features as the ConvGRU. The baseline MLP maps these features through 16- and 8-unit hidden layers, each with layer normalization and a SiLU activation, followed by a linear layer that produces logits for the three behavior classes.

3.8 Transport Model Parameter Tuning

We tuned the transport model on steady samples from the validation split to assess its ability to predict NO_2 plume evolution during periods of relatively stable NO_x emissions. We sampled 2,500 records and randomly selected one consecutive NO_2 raster transition from each record.

We evaluated the wind-blending factor λ, wind-speed scale γ, dispersion width σ_0, and decay timescale τ using the L1 loss between the transport model's predicted NO_2 raster $\widehat{N}_{i+1}$ and the observed raster N_{i+1}. The parameter set $\lambda = 0.5$, $\gamma = 0.8$, $\sigma_0 = 1$, and $\tau = 8$ yielded the lowest error from this procedure.

3.9 Training Procedure

We standardize the NO_2, wind, temperature, and transport residual rasters and the tabular inputs in every split using means and standard deviations computed from the training split.

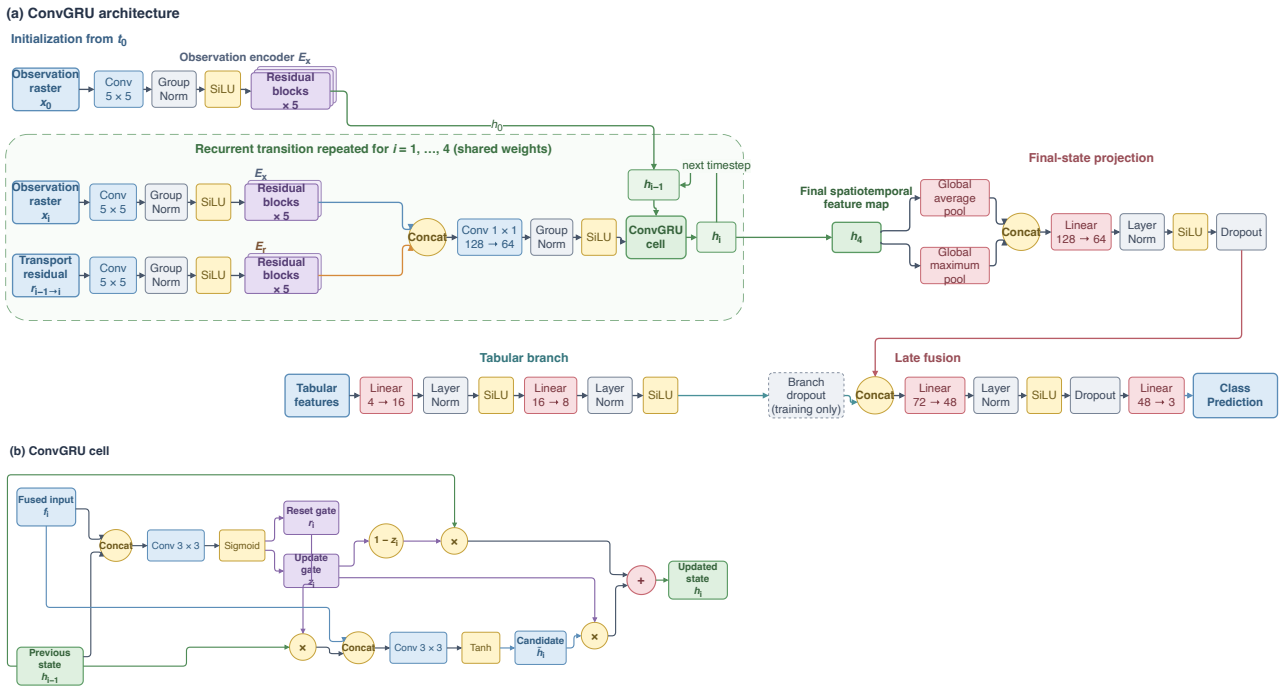


Figure 5: ConvGRU architecture for emissions change classification. (a) Observation and transport-residual encoders, four recurrent transitions, and late fusion with time-derived tabular features. (b) Reset, update, and hidden-state operations within one ConvGRU cell.

Table 2: Training parameters and setup details for the baseline MLP and ConvGRU. Entries spanning both model columns apply to both models; dashes indicate parameters not used by the MLP.

Parameter	Baseline MLP	ConvGRU
Trainable parameters	291	724,939
Model width	–	64
Optimizer	AdamW	
Batch size	128	128
Initial learning rate	10^{-3}	6.04×10^{-5}
Weight decay	10^{-4}	10^{-4}
Activation dropout	–	0.20
Tabular branch dropout	–	0.25
Gradient-norm limit	5.0	
LR scheduler	Reduce on plateau	
LR reduction factor	0.5	
Scheduler patience	5 epochs	
Selection metric	Validation macro-F1	
Early-stopping patience	10 epochs	
Hardware	1 NVIDIA RTX A5000 GPU	

We tuned the ConvGRU by varying model width, initial learning rate, weight decay, and activation dropout, selecting the configuration with the highest validation macro-F1. Table 2 reports the selected values and other settings used for final model training.

For the ConvGRU model, we also apply tabular branch dropout during training: for each sample, the tabular embedding is fully zeroed with probability 0.25 before late fusion. This discourages the classifier from relying on the time-derived tabular features alone.

Both classifiers use inverse-frequency class weights computed from the final training split and minimize weighted cross-entropy.

4 Results

4.1 Classification Performance

We use F1 to summarize the balance between precision and recall for each behavior class, and macro-F1 to give the three classes equal weight despite unequal sample counts. Figure 6 compares the models on the test set. The ConvGRU achieves a macro-F1 of 0.56, compared with 0.45 for the baseline MLP.

ConvGRU macro-F1 was 0.62 on training and 0.60 on validation, while the baseline MLP scored 0.53 on both. Both models scored lower on the test set, but the ConvGRU's advantage persisted.

The ConvGRU's largest test set F1 gain is for steady samples (0.142 points). The monthly class composition in Figure 7 suggests why: increases are more common in summer and decreases in winter, while steady samples show a weaker seasonal pattern, giving the time-only baseline less information. Decrease remains the ConvGRU's weakest class, with test recall of 0.471 versus 0.527 for the baseline MLP. This may reflect fewer decrease training samples and a delayed raster response: NO_2 emitted before a decline remains visible as the plume advects, disperses, and decays.

4.2 Performance Across Evaluation Strata

4.2.1 Performance Across Months. Figure 7 shows that ConvGRU accuracy varies less by month than baseline MLP accuracy. Their

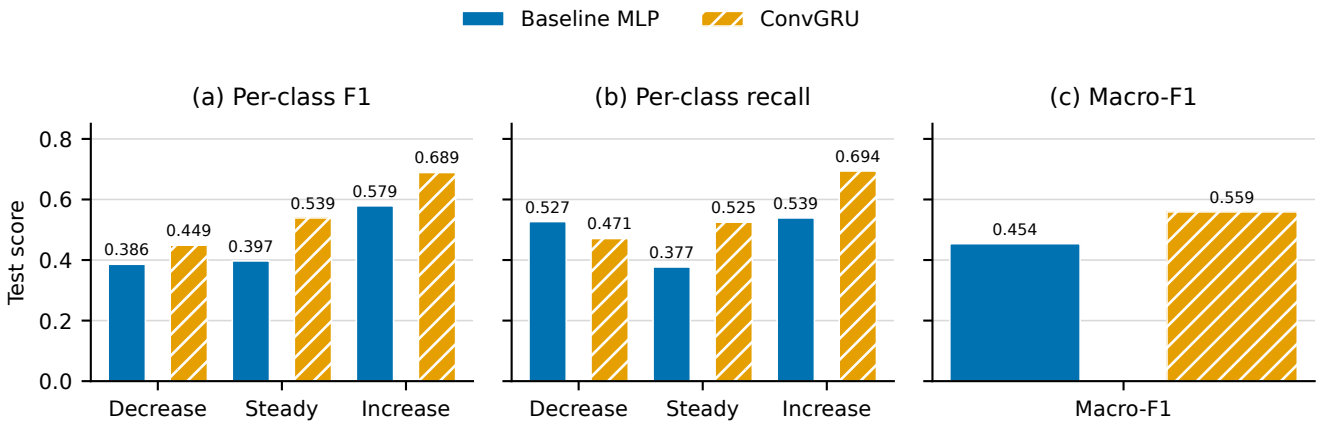


Figure 6: Test set performance of the baseline MLP and ConvGRU for emissions change classification. (a) F1 and (b) recall for the decrease, steady, and increase classes; (c) macro-F1 across classes.

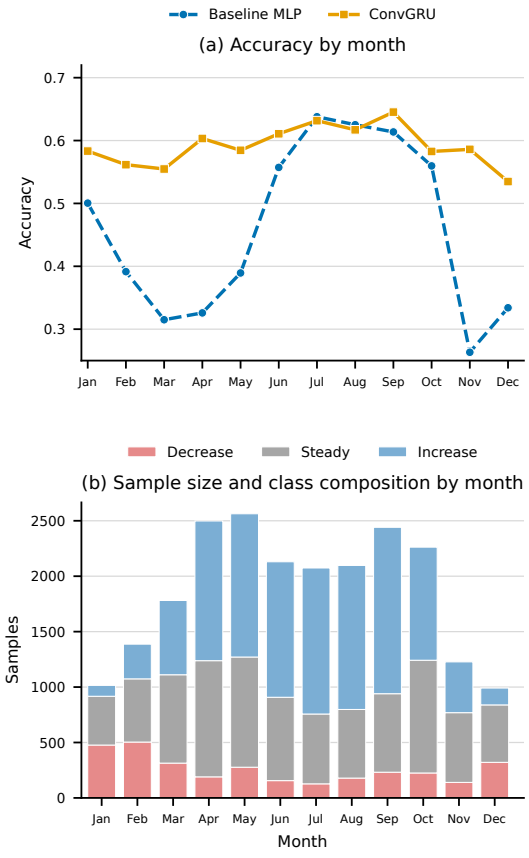


Figure 7: Test set performance and class composition by month. (a) Classification accuracy of the baseline MLP and ConvGRU. (b) Number of samples in each emissions change class (decrease, steady, and increase) by month.

results are closest in summer, but ConvGRU accuracy exceeds MLP accuracy by 0.277 in April and 0.323 in November. Panel (b) shows that increases are especially common in summer, while steady samples make up a larger share in spring and late fall, when the models differ more. This pattern is consistent with the MLP relying on seasonal class prevalence, whereas the ConvGRU can also use raster inputs to identify plant activity.

4.2.2 Performance by Local Solar Hour. Figure 8 shows that both models generally lose accuracy after the morning hours. The ConvGRU outperforms the baseline MLP from 08:00 through 16:00. Panel (b) shows that sample counts peak around 10:00–11:00 and are sparse at the earliest and latest hours; the class mix also changes through the day. These aggregate results therefore do not isolate an effect of solar hour from differences in class composition.

4.2.3 Performance by Census Region. Figure 9 shows ConvGRU macro-F1 of 0.57 in the South, 0.51 in the West, and 0.46 in the Midwest, while the baseline MLP remains near 0.45 in all three regions. The ConvGRU's regional pattern broadly tracks the class mix in panel (b): its score is highest where increase samples are most prevalent and lowest where decrease samples have the largest share. The West, where steady samples dominate, falls between them. Class composition may contribute to these differences, though this comparison does not isolate its effect. The test split contains no Northeastern AOIs.

4.2.4 Performance by NO_x Change Magnitude. Figure 10 shows that lower-magnitude decreases are the ConvGRU's weakest group among those plotted: its score is 0.59, compared with 0.69 for larger decreases. The baseline MLP scores 0.68 on smaller decreases, while the models perform similarly on larger decreases. By contrast, the ConvGRU's scores for smaller and larger increases are nearly identical (0.81 and 0.83). Smaller decreases may therefore be a principal source of its weaker performance on decrease events overall. This pattern is consistent with a lag in the raster response: NO_2 emitted before a unit's emissions decline or stop can remain visible as it advects and decays, making smaller decreases harder to distinguish.

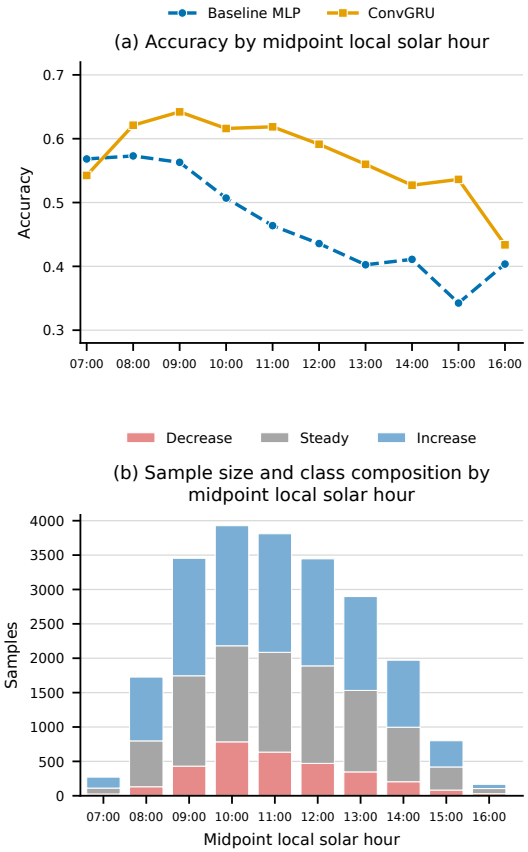


Figure 8: Test set performance and class composition by local solar hour at the sequence midpoint. (a) Classification accuracy of the baseline MLP and ConvGRU. (b) Number of samples in each emissions change class (decrease, steady, and increase) by hour.

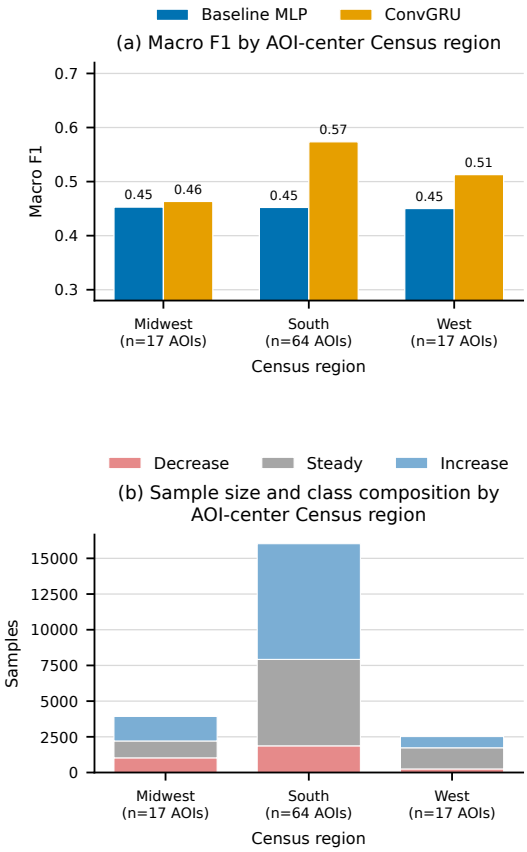


Figure 9: Test set performance and class composition by Census region of the AOI center. (a) macro-F1 of the baseline MLP and ConvGRU. (b) Number of samples in each emissions change class (decrease, steady, and increase) by region.

4.3 Ablation Study

We retrain the ConvGRU for each ablation in Table 3, zeroing the listed input for every sample. For timestep ablations, we zero the NO_2 and weather rasters but retain overlapping transport residuals. We omit t_0 because it initializes the hidden state, making its removal different from removing later observations.

Removing transport residuals causes the largest drop in test macro-F1 (12.58 percentage points), including a fall in steady F1 from 0.52 to 0.31. This suggests the model uses deviations from predicted plume transport to distinguish emissions changes from plume transport. Removing temperature costs 4.87 points and particularly hurts increase F1 (0.69 to 0.57). Temperature may help the model interpret atmospheric effects in the rasters, although it could also act as a proxy for electricity demand. Removing tabular features costs only 0.74 points, supporting the view that the ConvGRU identifies emissions trajectories primarily from visual inputs rather than seasonal timing alone. Removing wind rasters costs only 0.27 points, but wind also enters the transport residual calculation, so this does not establish that wind is uninformative.

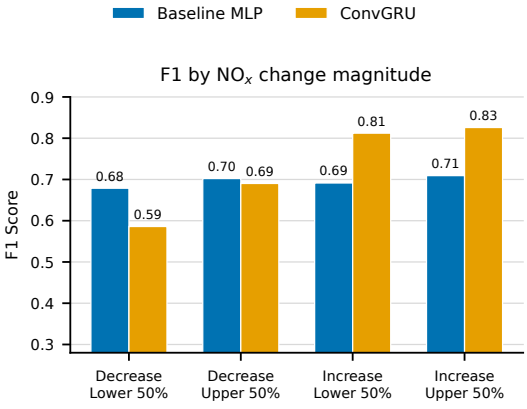


Figure 10: Test set performance by NO_x emissions change magnitude within increase and decrease classes. Within each class, samples are split into lower and upper halves at the median absolute emissions change.

Table 3: ConvGRU ablation performance on the test set. Each ablated model was retrained with the listed input zeroed out for all samples. Baseline is the full-input ConvGRU.

Input feature removed	macro-F1	Decrease F1	Steady F1	Increase F1
Baseline	0.56	0.46	0.52	0.69
Wind rasters	0.56	0.46	0.54	0.67
Temperature rasters	0.51	0.42	0.53	0.57
Transport residuals	0.43	0.36	0.31	0.62
Tabular features	0.55	0.46	0.51	0.68
t_1 observation	0.54	0.44	0.56	0.62
t_2 observation	0.54	0.47	0.54	0.62
t_3 observation	0.53	0.45	0.50	0.63
t_4 observation	0.54	0.47	0.51	0.64

5 Discussion

5.1 Evaluation of Model Results

The ConvGRU uses raster information about power plant activity beyond time of day and seasonality. Its advantage over the baseline MLP is greatest for steady samples, where features derived from time are less informative. Transport residuals contribute most to the model's performance in the ablation study, particularly for steady samples. This suggests that accounting for plume transport helps the model distinguish emissions behavior from changes in atmospheric conditions. The ConvGRU performs worse on decreases than on increases, possibly because decreases are less common in the training set and previously emitted NO_2 can remain visible after emissions fall.

Across the monthly, hourly, and regional subsets, both models tend to perform better when increases are more prevalent than when decreases are more prevalent. The ConvGRU's advantage tends to grow as the share of steady samples rises. A binary task that separates increases from all other behavior might therefore be easier to learn, since it would not require the model to distinguish decreases from steady periods. Future work could test this framing during operating hours when increasing activity is more common and its visual signal may be clearer.

5.2 Implications for Power Plant Monitoring

Satellite plume inversion depends on observations that distinguish a plant's plume from nearby emissions. Low source emissions, nearby sources, cloud screening, and very low or high wind speeds can reduce the number of usable observations or increase uncertainty in emission-rate estimates [12, 14, 20]. Our model instead classifies relative emissions changes across 72 km × 72 km AOIs while accounting for transport dynamics. This design could extend activity monitoring to some clustered or urban settings where plume inversion is difficult, although we have not directly evaluated that comparison.

This modeling approach may also be useful for detecting when specific units or entire facilities go offline or come back online due to either a planned outage or an equipment failure. Public access to unit-level outage information varies by independent system operator (ISO): CAISO publishes daily unit-level reports, ERCOT publishes unplanned resource outages after a three-day delay, and PJM releases only aggregate outage information [5, 7, 19]. For facilities subject to CEMS reporting, periods of nonoperation and subsequent restarts can instead be inferred retrospectively from quarterly EPA operating and emissions data [31, 34]. The current ConvGRU infers emissions trajectories across each AOI. A model that jointly predicts the activity label and a spatial contribution raster could identify which pixels support a predicted change and help isolate facilities whose units exhibit behavior consistent with an outage.

Another relevant activity is a substantial increase or decrease in a large unit's emissions that lasts several hours or more. Such changes can reflect routine dispatch rather than an outage, and they occur more frequently than outage events. Hourly generation by balancing authority and fuel is public, but unit-level dispatch or output may remain confidential or appear only after a delay, as in PJM and ERCOT [6, 19, 25]. Joint predictions of activity and pixel contributions could help identify which facilities in an AOI exhibit emissions changes consistent with a major increase or decrease in generation from one or more units.

Growing demand for power plants dedicated to data centers creates a need for remote sensing methods to monitor activity at facilities outside federal CEMS reporting programs [11, 27, 32]. Existing facilities already use multiple natural gas turbines to supply data centers, while planned facilities may use turbines or reciprocating engines to meet similarly large, relatively steady electricity loads [11, 26, 32]. Our model may offer a way to detect short-term changes in activity at these facilities. These activity estimates could complement inverse modeling estimates of emission rates, which require suitable observations and assumptions about atmospheric transport [14, 20].

5.3 Limitations and Future Work

The available record limits the size and coverage of the modeling dataset. Our observations span August 2023 through June 2026, which covers the full range of dates with both TEMPO observations and quarterly CEMS data available at the time of study. Cloud cover and pixel validity requirements further reduced the number of usable five-observation sequences. As the TEMPO and EPA CEMS records grow, more sequences may meet our sampling requirements. Beyond North America, GEMS observations could be paired with South Korea's publicly released stack NO_x measurements, while Sentinel-4 observations could be paired with European generation data for large individual units [8, 9, 13, 36]. These extensions would require new label construction and validation because the available ground data differ from EPA CEMS. Improved NO_2 retrieval and cloud-correction algorithms may also make changes in plume structure more discernible across our raster sequences [21].

Transport residuals produced the largest gain in our ablation study, but their calculation uses one effective wind field, simplified horizontal dispersion, and a fixed decay timescale. A future model could interpolate wind vector data based on stack height, account for boundary-layer mixing and plume rise, and model NO_2 chemistry as it varies with sunlight and temperature [12, 16]. Better predictions of how existing plumes move and persist could

reduce transport-related variation in the residuals, making emissions changes easier to detect, especially decreases obscured by lingering NO_2.

Our labels aggregate emissions from all facilities within each AOI. When nearby units change activity in opposing directions, their signals can partly cancel, leaving little visual evidence of the aggregate label. The current model cannot identify which facility drove a predicted change or what caused it. Future work could jointly predict the activity label and a per-pixel spatial contribution map to help associate the observed signal with specific facilities in an AOI.

6 Conclusion

We frame short-term power plant monitoring as a change detection task over sequences of TEMPO NO_2 observations rather than estimating an exact emission rate from each scan. On more than 22,000 held-out samples covering 116 active U.S. plants, the ConvGRU achieved a macro-F1 of 0.56, compared with 0.45 for a baseline using only time-derived features. This gain suggests that the raster sequences contain information beyond recurring daily and seasonal patterns, even though retrieval uncertainty, changing atmospheric conditions, and nearby sources can obscure individual plumes. Transport residuals compare each scan with a prediction of how an existing plume should move and decay, and our ablation study identified them as the most important model input. These findings suggest that change detection may be a more tractable framework for modeling short-term emissions from noisy satellite observations than direct regression of absolute emissions. This approach could support independent monitoring of plant activity during CEMS failures, before quarterly CEMS data are released publicly, and at facilities not subject to hourly reporting requirements.

References

[1] Ahmed Alnaim, Ziheng Sun, and Daniel Tong. 2022. Evaluating Machine Learning and Remote Sensing in Monitoring NO_2 Emission of Power Plants. *Remote Sensing* 14, 3 (2022), 729. doi:10.3390/rs14030729

[2] Nicolas Ballas, Li Yao, Chris Pal, and Aaron Courville. 2015. Delving Deeper into Convolutional Networks for Learning Video Representations. arXiv:1511.06432 [cs.LG] doi:10.48550/arXiv.1511.06432

[3] Samuel James Beaudry, Qindan Zhu, and Ronald C. Cohen. 2026. Quantification of Power Plant NO_x Emissions: A Reference for Diurnal and Seasonal Accuracy of TEMPO. doi:10.31223/X5B808 EarthArXiv preprint.

[4] Gedas Bertasius, Heng Wang, and Lorenzo Torresani. 2021. Is Space-Time Attention All You Need for Video Understanding?. In *Proceedings of the 38th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 139)*. PMLR, Virtual Event, 813–824.

[5] California Independent System Operator. 2026. Curtailed and Non-Operational Generators. https://www.caiso.com/market-operations/outages/curtailed-and-non-operational-generators Accessed October 19, 2026.

[6] Electric Reliability Council of Texas. 2026. 60-Day SCED Disclosure Reports. https://www.ercot.com/mp/data-products/data-product-details?id=NP3-965-ER Accessed October 19, 2026.

[7] Electric Reliability Council of Texas. 2026. Unplanned Resource Outages Report. https://www.ercot.com/mp/data-products/data-product-details?id=NP1-346-ER Accessed October 19, 2026.

[8] European Network of Transmission System Operators for Electricity. 2014. Transparency Platform: Detailed Data Descriptions. https://www.entsoe.eu/fileadmin/user_upload/_library/resources/Transparency/MoP%20Ref02%20-%20EMFIP-Detailed%20Data%20Descriptions%20V1R4-2014-02-24.pdf

[9] European Space Agency. 2025. Sentinel-4 Offers First Glimpses of Air Pollutants. https://www.esa.int/Applications/Observing_the_Earth/Copernicus/Sentinel-4/Sentinel-4_offers_first_glimpses_of_air_pollutants

[10] D. P. Finch, P. I. Palmer, and T. Zhang. 2022. Automated Detection of Atmospheric NO_2 Plumes from Satellite Data: A Tool to Help Infer Anthropogenic Combustion Emissions. *Atmospheric Measurement Techniques* 15 (2022), 721–733. doi:10.5194/amt-15-721-2022

[11] Kevin D. Gauld, Daniel J. Varon, Nicholas Balasus, and Daniel H. Cusworth. 2026. Quantifying AI Data Center Nitrogen Oxide (NO_x) Emissions from Space. arXiv:2608.22153 [physics.ao-ph] doi:10.48550/arXiv.2608.22153

[12] Ruizhe Huang and Sherrie Wang. 2026. Global Variability in the Detectability of Power Plant NO_2 Plumes from Space. *Atmospheric Measurement Techniques* 19 (2026), 6099–6124. doi:10.5194/amt-19-6099-2026

[13] Korea Environment Corporation. 2026. CleanSYS Stack Monitoring Results: Public Real-Time Data. https://www.data.go.kr/data/15136837/fileData.do Accessed October 19, 2026.

[14] Gerrit Kuhlmann, Erik Franciscus Maria Koene, Chloe Natasha Schooling, Paul Ian Palmer, Òscar Collado López, and Marc Guevara. 2026. Temporal Variability of NO_x Emissions from Power Plants: A Comparison of Satellite- and Inventory-Based Estimates. *Atmospheric Chemistry and Physics* 26 (2026), 4405–4421. doi:10.5194/acp-26-4405-2026

[15] Pieternel F. Levelt, Gijsbertus H. J. van den Oord, Marcel R. Dobber, Anssi Mälkki, Huib Visser, Johan de Vries, Piet Stammes, Jens O. V. Lundell, and Heikki Saari. 2006. The Ozone Monitoring Instrument. *IEEE Transactions on Geoscience and Remote Sensing* 44, 5 (2006), 1093–1101. doi:10.1109/TGRS.2006.872333

[16] Samuel Meier, E. F. M. Koene, Maarten Krol, Dominik Brunner, Alexander Damm, and Gerrit Kuhlmann. 2024. A Lightweight NO_2-to-NO_x Conversion Model for Quantifying NO_x Emissions of Point Sources from NO_2 Satellite Observations. *Atmospheric Chemistry and Physics* 24 (2024), 7667–7686. doi:10.5194/acp-24-7667-2024

[17] National Oceanic and Atmospheric Administration. 2026. High-Resolution Rapid Refresh (HRRR). https://rapidrefresh.noaa.gov/hrrr/ Accessed October 15, 2026.

[18] National Oceanic and Atmospheric Administration. 2026. NOAA High-Resolution Rapid Refresh (HRRR) Model. https://registry.opendata.aws/noaa-hrrr-pds/ Accessed October 15, 2026.

[19] PJM Interconnection. 2026. Data Availability. https://www.pjm.com/markets-and-operations/etools/data-miner-2/data-availability.aspx Accessed October 18, 2026.

[20] D. Santaren, J. Hakkarainen, G. Kuhlmann, E. Koene, F. Chevallier, I. Ialongo, H. Lindqvist, J. Nurmela, J. Tamminen, L. Amorós, D. Brunner, and G. Broquet. 2025. Benchmarking Data-Driven Inversion Methods for the Estimation of Local CO_2 Emissions from Synthetic Satellite Images of XCO_2 and NO_2. *Atmospheric Measurement Techniques* 18 (2025), 211–239. doi:10.5194/amt-18-211-2025

[21] S. Seo, P. Valks, R. Lutz, K.-P. Heue, P. Hedelt, V. Molina García, D. Loyola, H. Lee, and J. Kim. 2024. Tropospheric NO_2 Retrieval Algorithm for Geostationary Satellite Instruments: Applications to GEMS. *Atmospheric Measurement Techniques* 17 (2024), 6163–6191. doi:10.5194/amt-17-6163-2024

[22] Xingjian Shi, Zhourong Chen, Hao Wang, Dit-Yan Yeung, Wai-Kin Wong, and Wang-Chun Woo. 2015. Convolutional LSTM Network: A Machine Learning Approach for Precipitation Nowcasting. In *Advances in Neural Information Processing Systems*, Vol. 28. Curran Associates, Inc., Montreal, Canada, 802–810.

[23] Kang Sun, Jobaer Ahmed Saju, Caroline R. Nowlan, Gonzalo González Abad, and Xiong Liu. 2025. Hourly Nitrogen Oxides Emissions Estimated From TEMPO and Comparison With Facility-Level Monitoring Data. *Journal of Geophysical Research: Atmospheres* 130, 23 (2025), e2025JD044565. doi:10.1029/2025JD044565

[24] TEMPO Science Team. 2026. *TEMPO Trace Gas and Cloud Level 2 and 3 Data Products: User Guide.* NASA Langley Atmospheric Science Data Center. doi:10.5067/doc/tempo/L2andL3UserGuide/V2.1 Version 2.1.

[25] U.S. Energy Information Administration. 2026. About the EIA-930 Data. https://www.eia.gov/electricity/930-content/about.html Accessed October 19, 2026.

[26] U.S. Energy Information Administration. 2026. Data Center Server Energy Use Grows Across the Commercial Building Stock. https://www.eia.gov/todayinenergy/detail.php?id=67704

[27] U.S. Energy Information Administration. 2026. Fossil Generation Could Rise with Faster-than-Expected Growth in Data Center Power Demand. https://www.eia.gov/todayinenergy/detail.php?id=67344

[28] U.S. Environmental Protection Agency. 2016. *Integrated Science Assessment for Oxides of Nitrogen—Health Criteria.* Technical Report EPA/600/R-15/068. U.S. Environmental Protection Agency. https://www.epa.gov/isa/integrated-science-assessment-isa-oxides-nitrogen

[29] U.S. Environmental Protection Agency. 2021. *Substitute Data in EPA CAMD's Power Sector Emissions Data.* Technical Report. U.S. Environmental Protection Agency. https://www.epa.gov/system/files/documents/2021-06/monitoring-insights-substitute-data.pdf

[30] U.S. Environmental Protection Agency. 2026. Air Pollutant Emissions Trends Data. https://www.epa.gov/air-emissions-inventories/air-pollutant-emissions-trends-data Accessed October 4, 2026.

[31] U.S. Environmental Protection Agency. 2026. CAMD's Power Sector Emissions Data Guide. https://nepis.epa.gov/Exe/ZyPURL.cgi?Dockey=P1019ZHX.txt Accessed October 19, 2026.

[32] U.S. Environmental Protection Agency. 2026. Clarification of Acid Rain Program Regulatory Provisions Concerning Islanded Power Generation Facilities.

https://www.epa.gov/system/files/documents/2026-07/epa-issues-clarification-on-islanded-power-generators-and-acid-rain-program-provisions.pdf

[33] U.S. Environmental Protection Agency. 2026. Emissions Monitoring and Reporting. https://www.epa.gov/power-sector/emissions-monitoring-and-reporting Accessed October 4, 2026.

[34] U.S. Environmental Protection Agency. 2026. Power Sector Data. Office of Atmospheric Protection, Clean Air and Power Division, Washington, DC. https://campd.epa.gov/ Accessed October 15, 2026.

[35] J. P. Veefkind, I. Aben, K. McMullan, H. Förster, J. de Vries, G. Otter, J. Claas, H. J. Eskes, J. F. de Haan, Q. Kleipool, M. van Weele, O. Hasekamp, R. Hoogeveen, J. Landgraf, R. Snel, P. Tol, P. Ingmann, R. Voors, B. Kruizinga, R. Vink, H. Visser, and P. F. Levelt. 2012. TROPOMI on the ESA Sentinel-5 Precursor: A GMES Mission for Global Observations of the Atmospheric Composition for Climate, Air Quality and Ozone Layer Applications. *Remote Sensing of Environment* 120 (2012), 70–83. doi:10.1016/j.rse.2011.09.027

[36] L. H. Yang, D. J. Jacob, R. Dang, Y. J. Oak, H. Lin, J. Kim, S. Zhai, N. K. Colombi, D. C. Pendergrass, E. Beaudry, V. Shah, X. Feng, R. M. Yantosca, H. Chong, J. Park, H. Lee, W.-J. Lee, S. Kim, E. Kim, K. R. Travis, J. H. Crawford, and H. Liao. 2024. Interpreting Geostationary Environment Monitoring Spectrometer (GEMS) Geostationary Satellite Observations of the Diurnal Variation in Nitrogen Dioxide (NO_2) over East Asia. *Atmospheric Chemistry and Physics* 24 (2024), 7027–7039. doi:10.5194/acp-24-7027-2024

[37] Peter Zoogman, Xiong Liu, R. M. Suleiman, W. F. Pennington, David E. Flittner, Jassim A. Al-Saadi, Brian B. Hilton, Dennis K. Nicks, Michael J. Newchurch, J. L. Carr, Scott J. Janz, M. R. Andraschko, Antti Arola, Barry D. Baker, Brian P. Canova, Christopher Chan Miller, Ronald C. Cohen, J. E. Davis, M. E. Dussault, David P. Edwards, Jack Fishman, A. Ghulam, Gonzalo González Abad, Michel Grutter, Jay R. Herman, J. Houck, Daniel J. Jacob, Joanna Joiner, Brian J. Kerridge, Jhoon Kim, Nickolay A. Krotkov, Lok N. Lamsal, Can Li, Anders Lindfors, Randall V. Martin, C. Thomas McElroy, Chris McLinden, Vijay Natraj, Daniel O. Neil, Caroline R. Nowlan, E. J. O'Sullivan, Paul I. Palmer, R. Bradley Pierce, Mark R. Pippin, Alfonso Saiz-Lopez, Robert J. D. Spurr, James J. Szykman, Omar Torres, J. Pepijn Veefkind, Ben Veihelmann, Hui Wang, Jun Wang, and Kelly Chance. 2017. Tropospheric Emissions: Monitoring of Pollution (TEMPO). *Journal of Quantitative Spectroscopy and Radiative Transfer* 186 (2017), 17–39. doi:10.1016/j.jqsrt.2016.05.008